

Lares: An Architecture for Secure Active Monitoring Using Virtualization

Bryan D. Payne, Martim Carbone,
Monirul Sharif, Wenke Lee

2008 IEEE Symposium on Security and Privacy

Motivation

- Host-based security tools are not adequately protected.
 - Antivirus
 - IDS
- We need *active* monitoring for *prevention*
- **Lares** - Applies VM techniques for active monitoring
- Presents an architecture and sample implementation

Active vs. Passive Monitoring

Passive Monitoring

- Run external scanning or polling
- Detect a malicious access (e.g. memory corruption) and Correct
- Example: VM Introspection

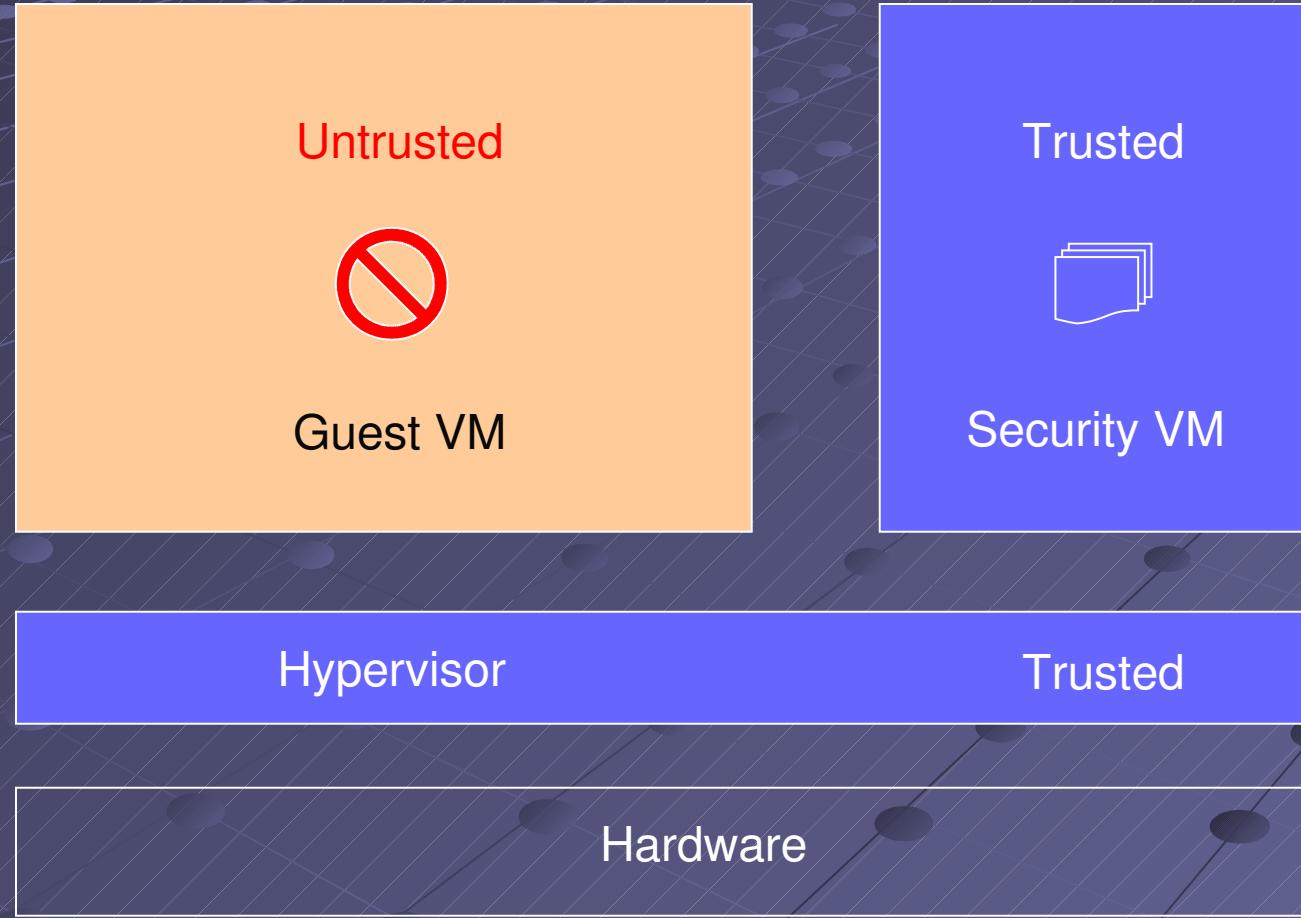
Active Monitoring

- Trigger a handler when an event happens
- Prevent a malicious access
- Example: Antivirus, IDS

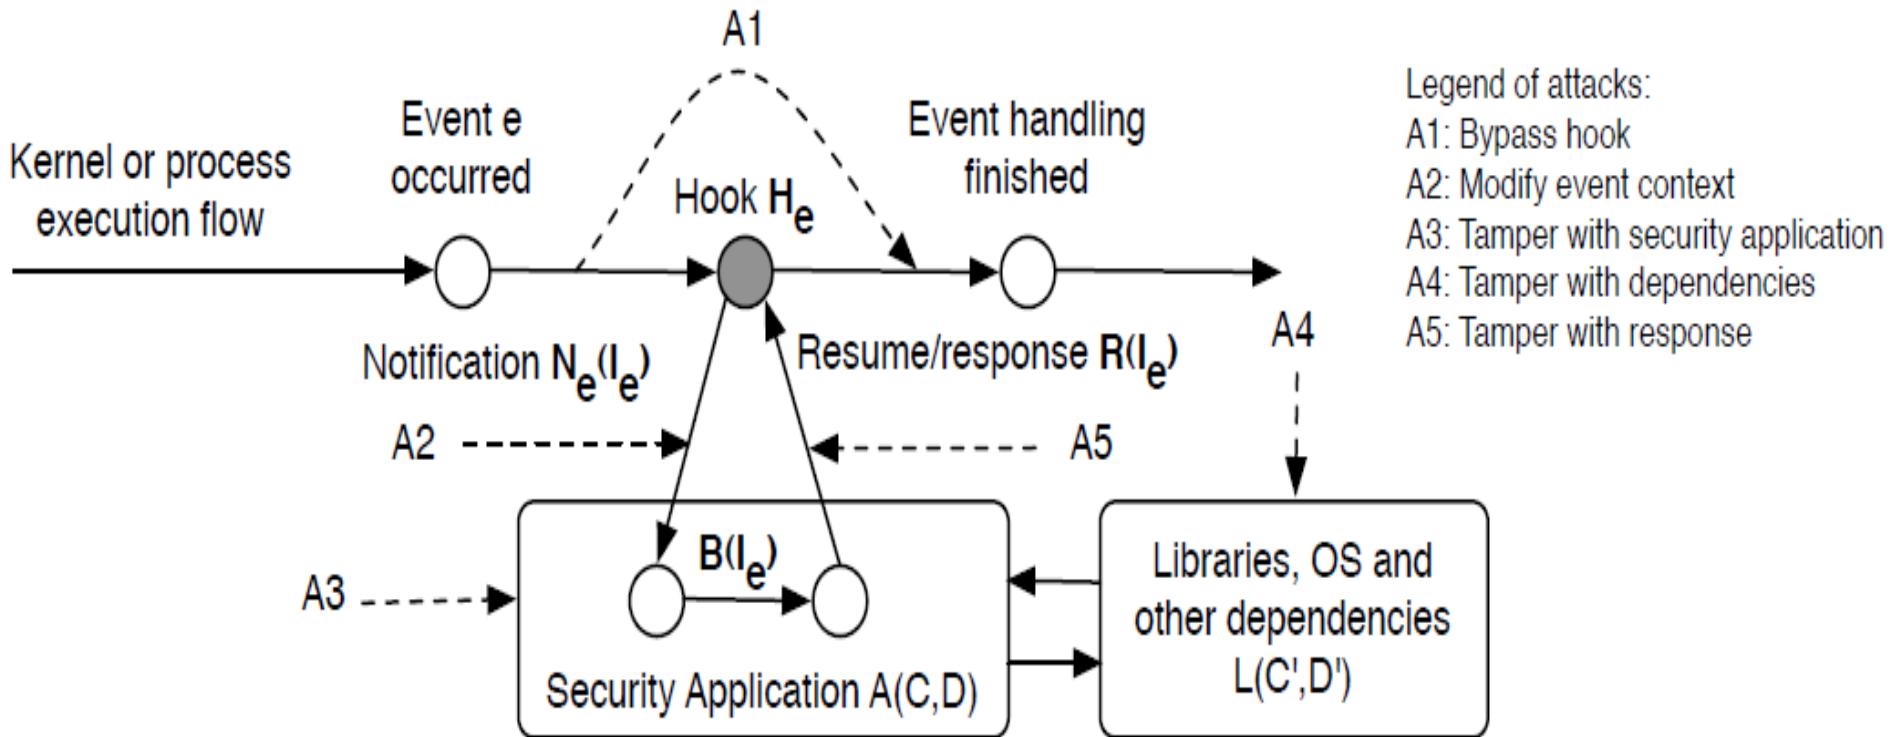
Major Challenges

- Malware monitoring requires
 - The analyzer and monitor must be hidden from malware
- Critical requirements:
 - For active monitoring we need to place hooks on the running system
 - Hooks are in the untrusted OS
 - How to protect the hooks and its comm. to VM?

System Layout



Adversary Model



- N_e is triggered if and only if e occurs legitimately
- I_e is not modifiable between occurrences of e and N_e
- 3. $B(I_e)$ of the security application is not altered maliciously
- 4. The effects of $R(I_e)$ on the system are enforced

Design Goals

- Protect Monitoring Components
 - Meet formal requirements
- Flexible Hook Placement
 - Difficult since adversary controls the host system
- Acceptable Performance Overhead

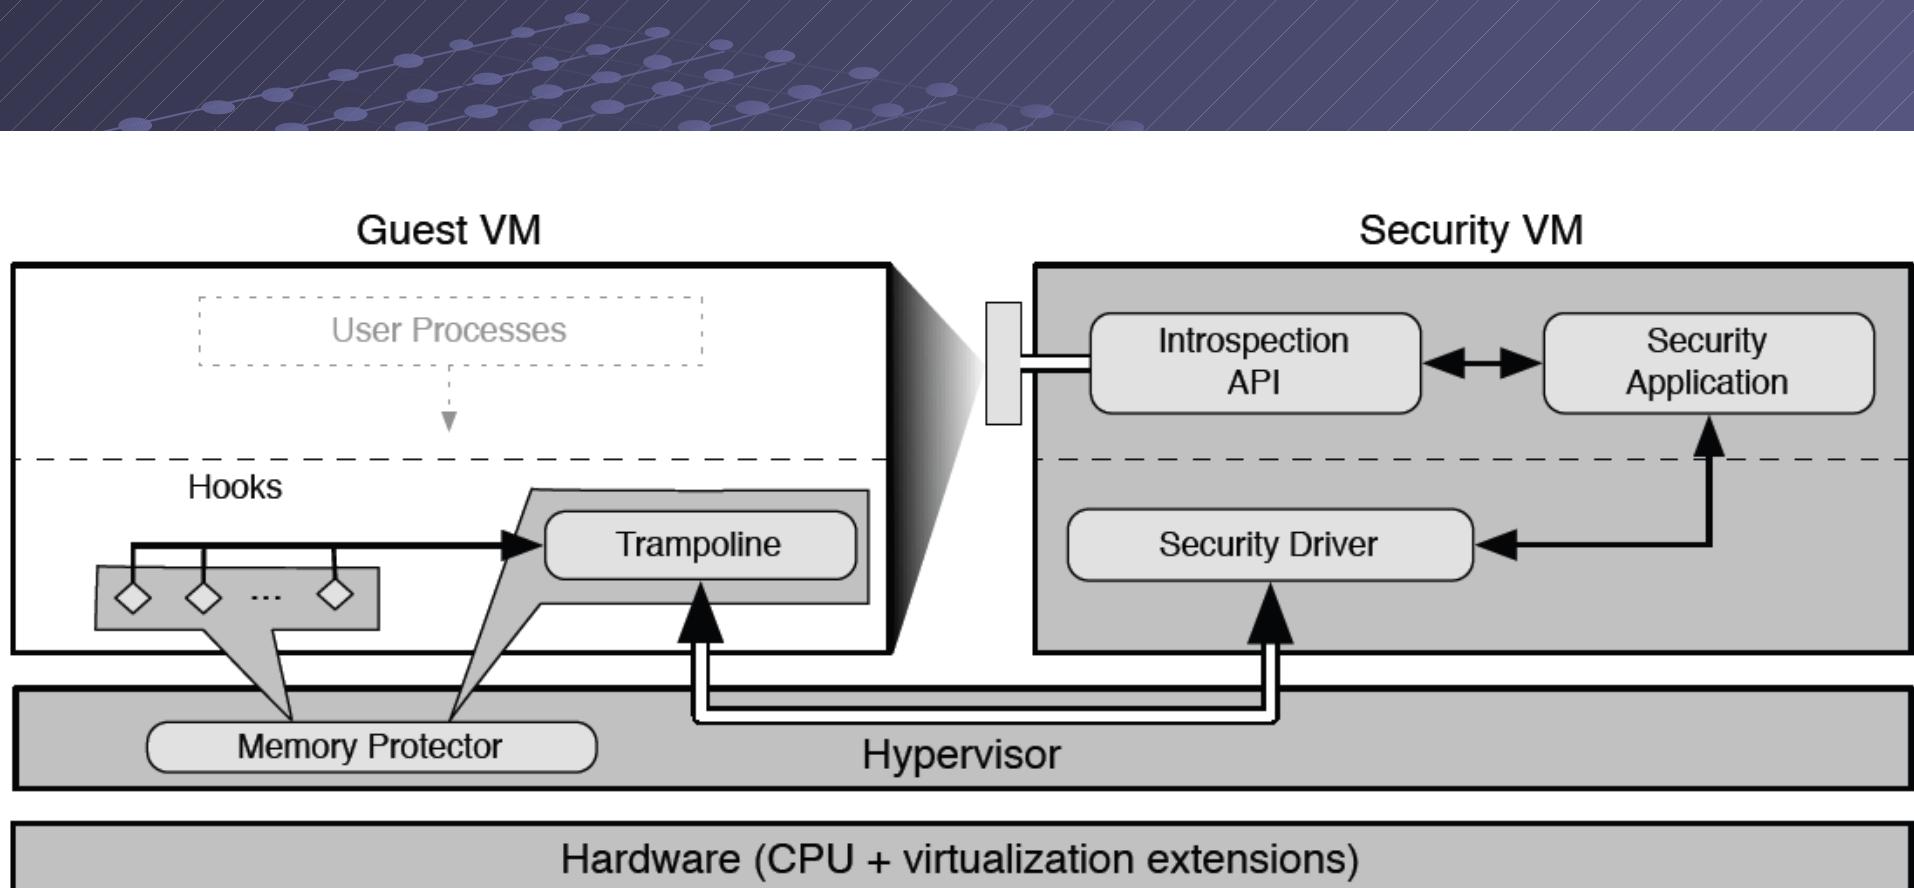
Solution Approach

- ➊ Split security solution over two VMs
- ➋ Guest VM
 - VM running the host OS
 - Controlled by adversary
 - Contains hooks
- ➌ Security VM
 - Contain security application (antivirus, IDS)
 - Trusted
 - Runs at same privilege level as hypervisor
- ➍ Use hypervisor-based memory protection to secure the hooks

Assumptions

- The hypervisor and Security VM form a trusted code base (TCB)
- The machine can undergo *secure boot*
- Guest VM undergoes an *initialization* after boot
 - Start the security components
 - Protect them
 - Add other security configurations

Architecture Overview



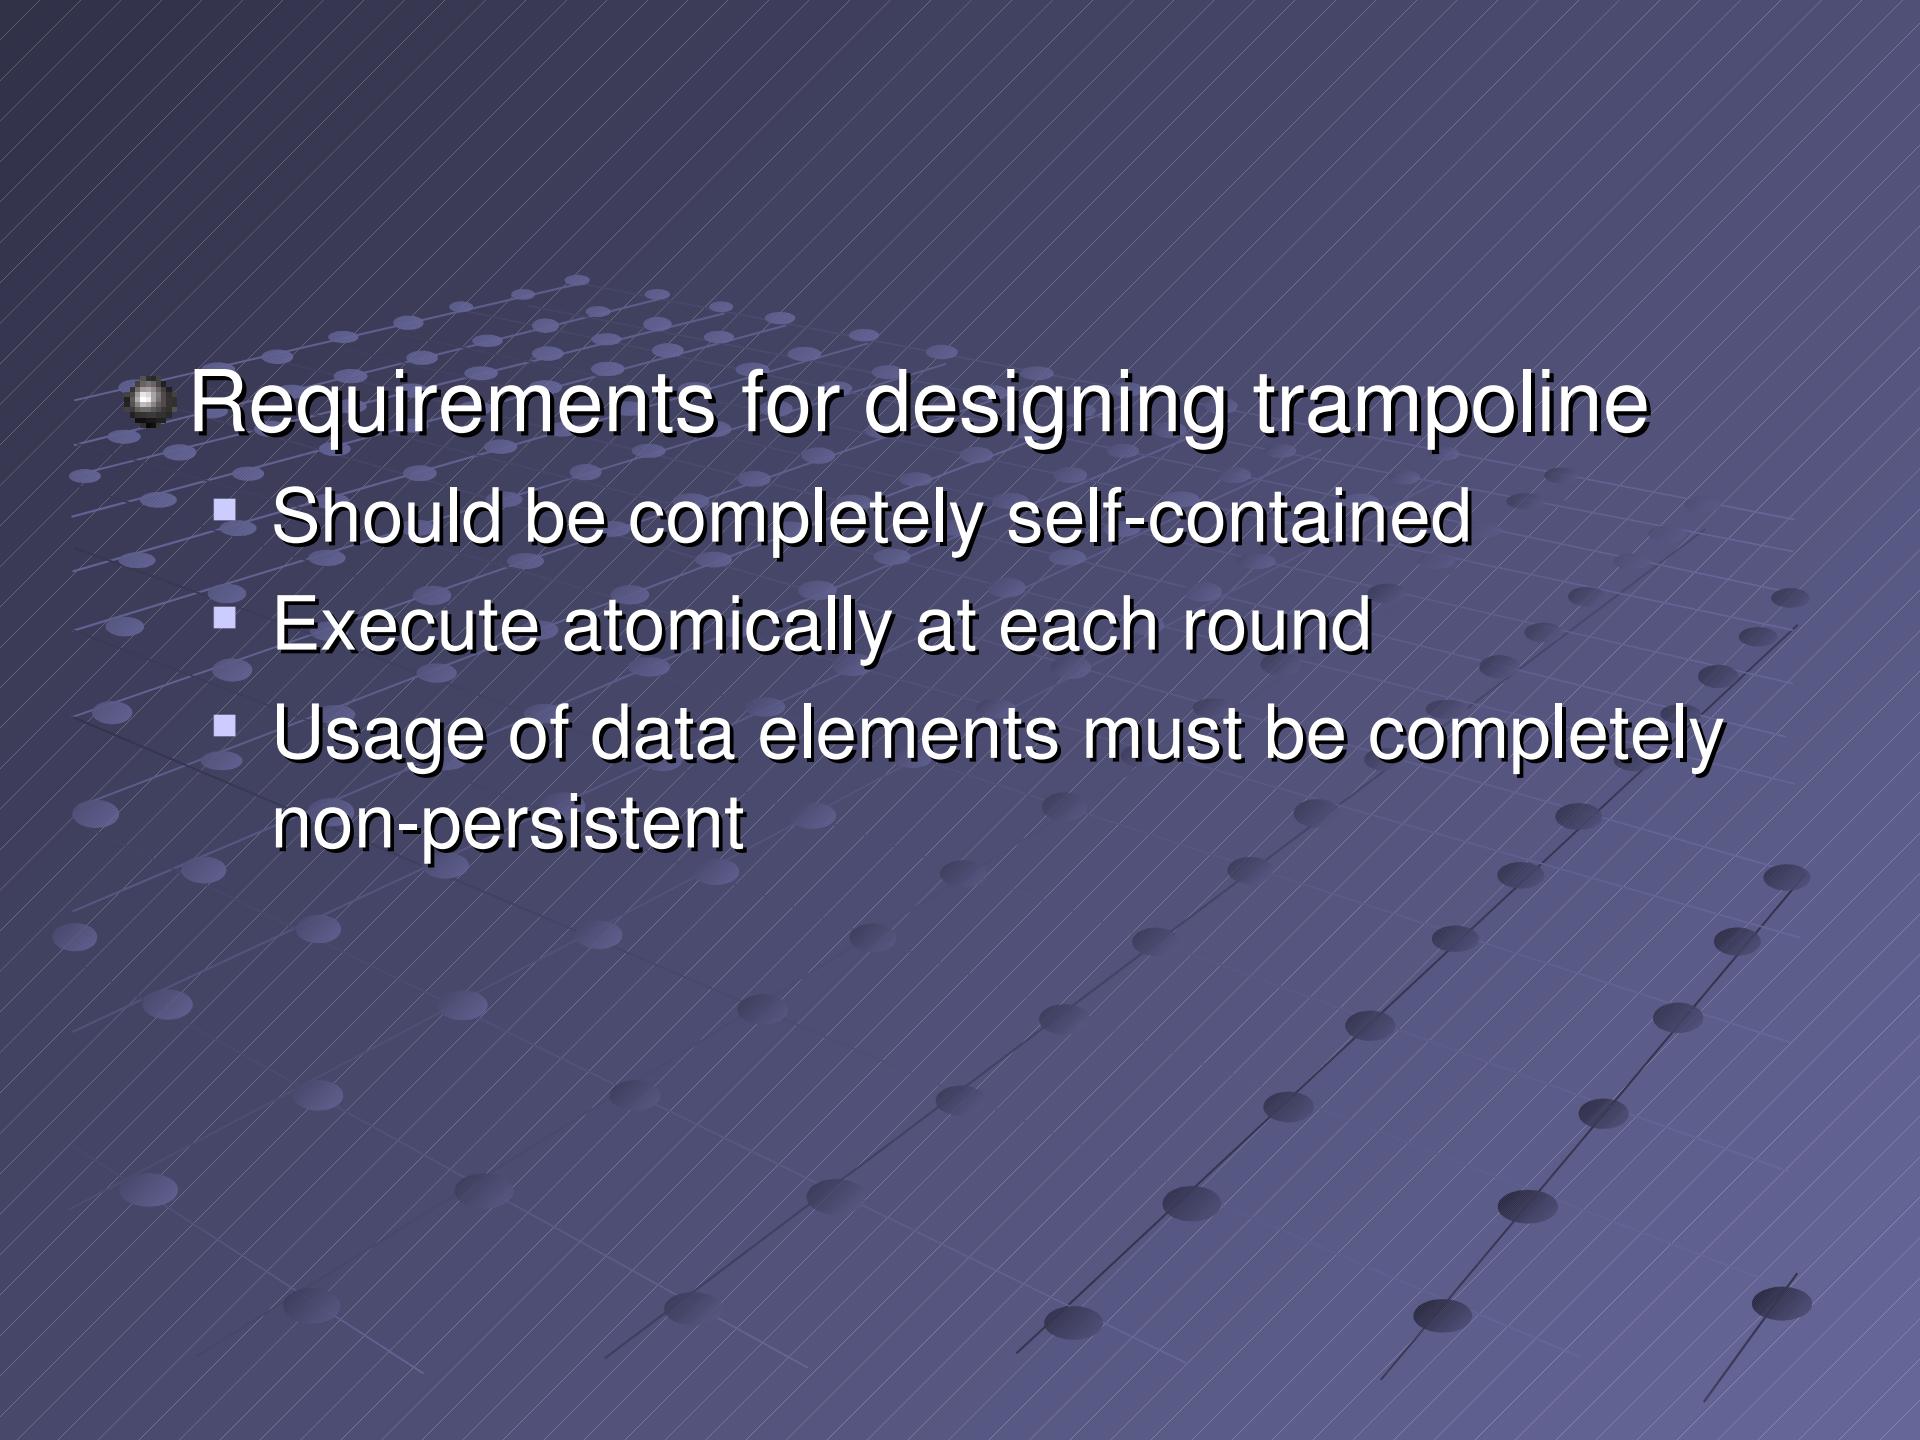
Guest VM Components

- Hooks

- JUMP in program code
- Redirection in jump tables
- Any other technique to transfer control

- Trampoline

- Bridge between the hooks and security drivers in security VM
- Passes arguments to the security driver



• Requirements for designing trampoline

- Should be completely self-contained
- Execute atomically at each round
- Usage of data elements must be completely non-persistent

Security VM Components

- Security application
- Introspection API
- Security driver

Hypervisor Components

- Guest OS component protection
 - Mark the memory location of hooks and trampoline as read-only
- Inter-VM communication
 - Delay returning to Guest OS until security VM gives a response

Implementation

- Uses hook on Windows syscall NtCreateSection
- Wrote a driver to install hook and trampoline
 - Initiated during Guest OS initialization
 - Never swapped to disk
 - Hypervisor activates memory protection
 - 324 SLOC

Continued ...

● Inter-VM Communication

- Added a new hypercall
- Uses a shared memory to send (receive) event request (response) to (from) hypervisor
- 127 SLOC

● Security Driver

- Imparts modularity/flexibility
- Interfaces with security application
- Built as Linux Kernel Module (LKM)
- 182 SLOC

Continued ...

- Security application

- Checks the file handle that was passed to NtCreateSection
- Allow execution if permitted
- 298 SLOC

Memory Protection

- Takes advantage of shadow paging in Xen
- Modify Xen's page fault handler to enable byte-sized granularity
- 78 SLOC

Guest page table

Virtual	Physical

PTE propagation

Shadow page table

Virtual	Machine

NO

Is the PF listed as protected?

YES

Mark as read-only

Page fault due to failed write

YES

Propagate exception to guest

NO

Is the write targeted at a protected region?

Emulate the write

Performance

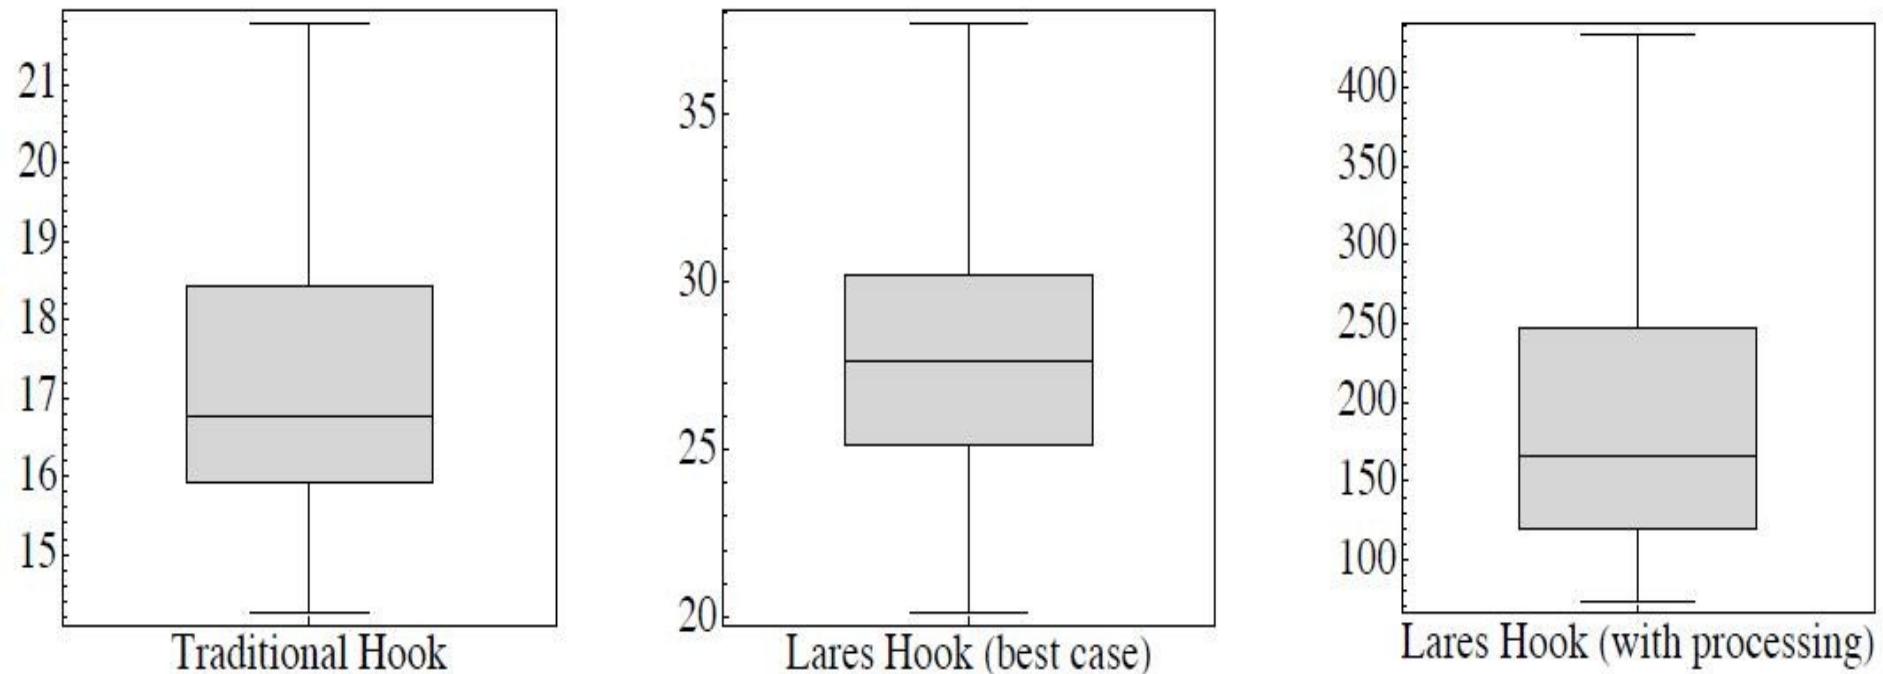
Settings:

- Intel Core Duo, 2GB RAM
- Hypervisor: Xen
- Guest VM: Windows XP, 384MB
- Security VM: Fedora

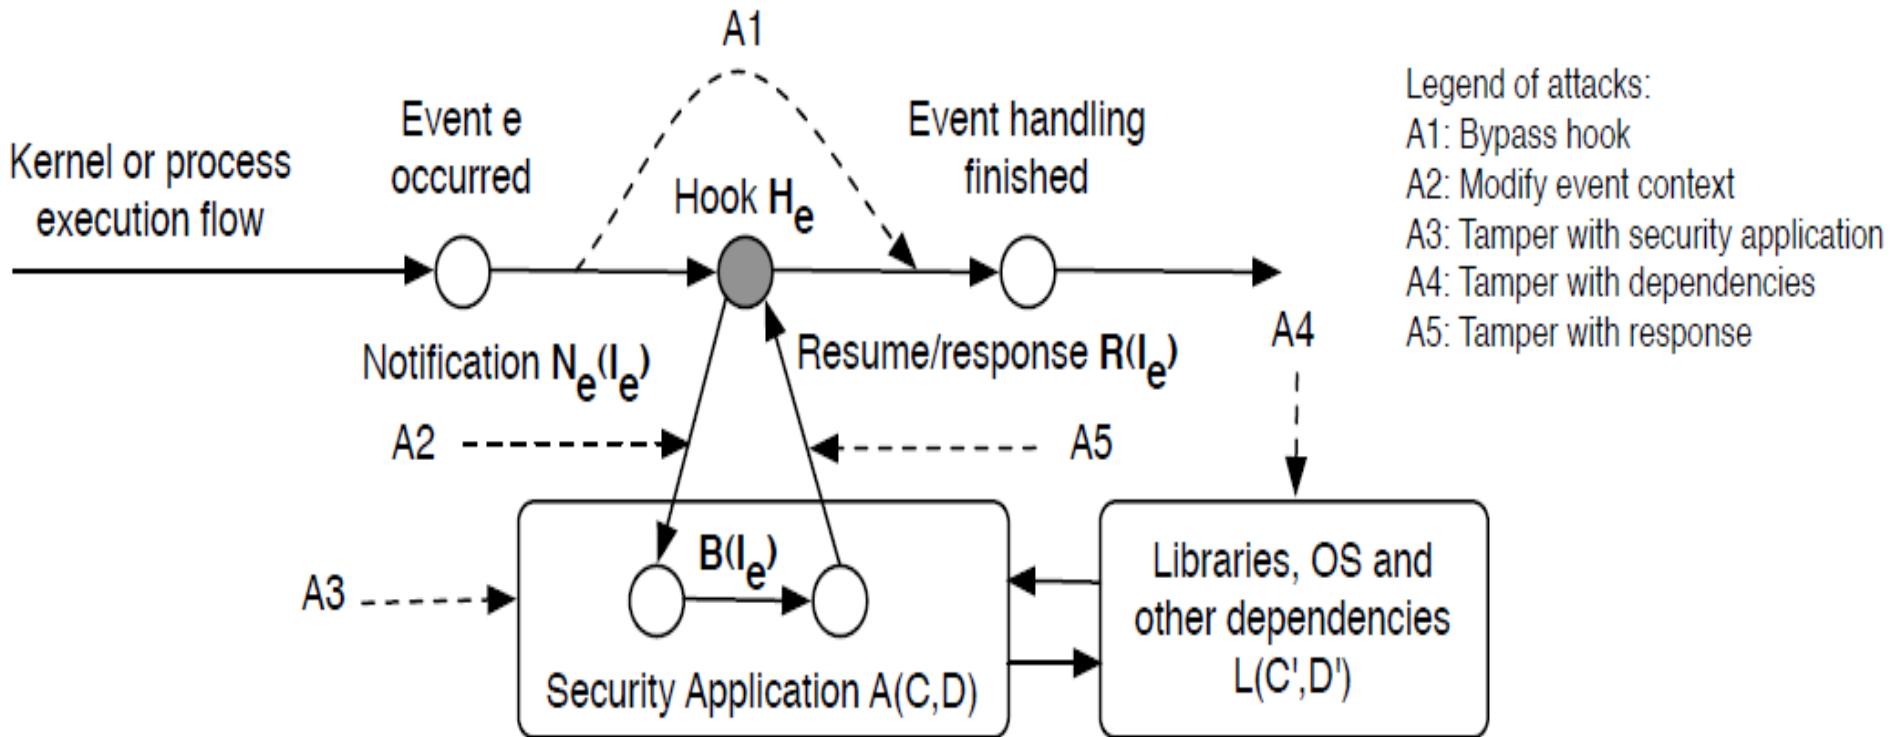
Hook processing time:

- Traditional Hook (security application in same VM as host)
 - 17 micro-sec
- Lares Hook (best case, minimal processing in security application)
 - 28 micro-sec
- Lares Hook (with processing)
 - 150 micro-sec

Results



Adversary Model

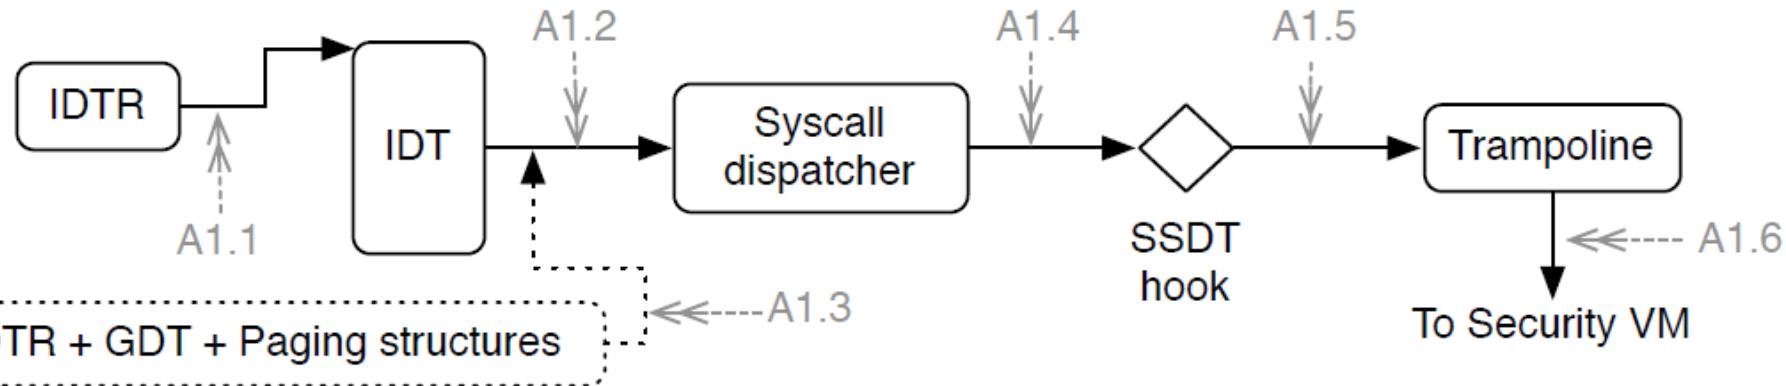


- N_e is triggered if and only if e occurs legitimately
- I_e is not modifiable between occurrences of e and N_e
- 3. $B(I_e)$ of the security application is not altered maliciously
- 4. The effects of $R(I_e)$ on the system are enforced

Security Analysis

- A3, A4 are met as per trusted VM assumptions
- A2, A5 are met by disabling interrupts by the trampoline
- A5 may be circumvented by non-maskable interrupts (NMI) [hardware fatal errors]
- But one VM uses a single core only
 - Malware cannot use the other core to send NMI

Handling Attacks of Type A1



- A1.5 and A1.6 handled by memory protection
- Write protect IDT to prevent A1.2 and A1.4
- A1.1 cannot be handled by Intel VT-x
- AMD SVM can trap changes to IDTR

Continued ...

- ➊ A1.3 is more difficult to handle

- Changes to GDTR, GDT, Page Tables
- Memory Introspection has similar problems
- Monitor shadow PT entries to check that they point to known good locations

- ➋ Preventing bogus event notification

- Mark the trampoline memory region as non-executable (NX)
- On fault, check origin of branch
 - ➌ If valid hook, proceed
 - ➌ Else attacker is making bogus call to trampoline

Conclusion

- Enable active monitoring with VM security
- Takes advantages of secure VM
- Can place hooks at arbitrary location in Kernel
- Low overhead
- Works on production systems
- Strong adversary model
- No special hardware needed